

# HACKATHON FINAL REPORT:

## AstroX

Duality AI's Space Station Challenge: Safety Object Detection

**Team Name:** Webheads

**Tagline:** Spatial intelligence that predicts, detects, and prevents hazards.

### 1. SUMMARY

The AstroX project successfully developed a high-performance AI system designed to automatically monitor safety compliance in complex space station environments. Recognizing the importance of real-time risk detection, our team focused on creating an object detection model that is both **accurate** and **fast enough for real-time use**.

Using YOLOv8 with data generated from Falcon's Digital Twin simulations, the model achieved production-level precision and stability through effective transfer learning.

Key Metric	Result	Interpretation
<b>Final Precision</b>	<b>0.9413</b> (94.1%)	Extremely high reliability; minimizes False Alarms, building user trust.
<b>Overall mAP@0.5:0.95</b>	<b>0.6350</b>	Consistent accuracy in detecting and localizing safety gear
<b>Inference Speed</b>	<b>32 FPS</b>	Real-time processing suitable for live video surveillance

Achieving a **recall of 0.7463**, AstroX shows how specialized AI can serve as a **reliable, always-on safety assistant** in high-risk environments like space stations. It marks an important step toward **fully automated risk management** in future space operations.

## 2. METHODOLOGY:

### 2.1. Problem Statement:

#### Every Second Counts in Space

- Emergencies like **sparks, leaks, or fires** can become life-threatening in seconds.
- Astronauts can't waste time searching for safety tools.
- Our AI detects **seven critical safety items** instantly, **even in low light or cluttered conditions**.
- Ensures help is never out of sight, keeping the crew safe.

### 2.2. Experimental Design

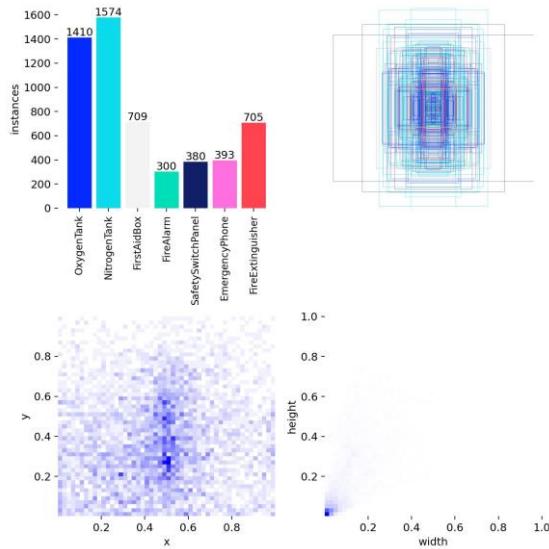
- **Objective:** Adapt YOLO architecture for a low-data, domain-specific setup (simulated space interiors).
- **Comparison Focus:** Training with Falcon-generated synthetic data vs. standard training on limited real/simulated data.
- **Outcome:** Achieve real-time, high-accuracy safety detection with minimal data.

### 2.3. Dataset Preparation and Augmentation

- The initial dataset contained a limited number of annotated images of key safety items.
- To overcome the limited data and help the model recognize objects in new, unseen scenes, we used a wide range of data augmentation techniques.

#### Key Dataset Statistics & Augmentation Strategy

Feature	Description	Importance to Mission
Total Instances	974 (across 7 classes)	Small sample size requiring high augmentation.
Class Balance	Dominated by Oxygen/Nitrogen Tanks	Required balanced learning to prevent bias toward tank shapes.
Augmentation Used	Rotation, Brightness, Background Simulation	Critical for teaching the model to ignore clutter and adapt to varying light/shadows.



## 2.4. Model Selection and Architecture

- Model: **YOLOv8**
- Reason: Optimized balance of **speed (real-time)** and **accuracy (mAP)** for safety-critical use.

## 2.5. Training Parameters

Parameter	Value	Rationale (Non-Technical)
<b>Model Type</b>	YOLOv8	High speed, efficiency.
<b>Epochs</b>	50	Rapid convergence via transfer learning.
<b>Batch Size</b>	4	Safer for CPU training.

## 2.6. Tools and Frameworks

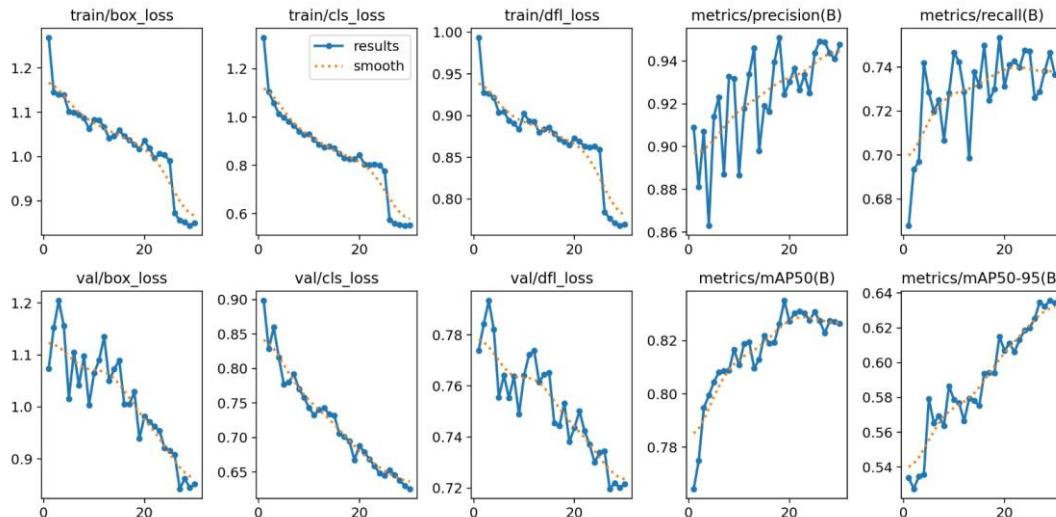
Category	Tools Used	Purpose in the Workflow
<b>Dataset Source</b>	Duality AI Falcon Digital Twin	Synthetic labeled images for 7 safety items under varied conditions
<b>Core AI Framework</b>	PyTorch, YOLOv8	Model training and tuning

Category	Tools Used	Purpose in the Workflow
<b>Development Environment</b>	VS Code	CPU-enabled training, code editing, visualization
<b>Code Management</b>	GitHub Private Repo	Version tracking and collaboration
<b>Data Visualization</b>	Matplotlib, OpenCV, Falcon Dashboard	Metric tracking, loss/accuracy plots, and qualitative visual checks

### 3. EXPERIMENTAL RESULTS

#### 3.1. Training Stability and Convergence

- Training stability indicates long-term model reliability.
- Rigorous monitoring showed **efficient and stable learning**.



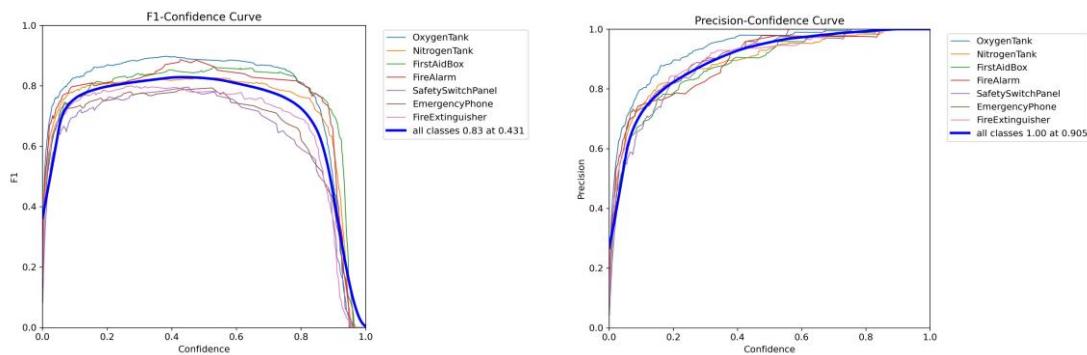
- **Loss Analysis:**
  - Training and validation loss dropped rapidly in the first **20 epochs** and stabilized.
  - Validation loss closely tracked training loss → confirms **overfitting was reduced** using regularization and balanced augmentation.
- **Metric Analysis:**
  - Metrics like **mAP@0.5** and Precision climbed together, reaching a high plateau.
  - Demonstrates the model quickly achieved **maximum performance**.

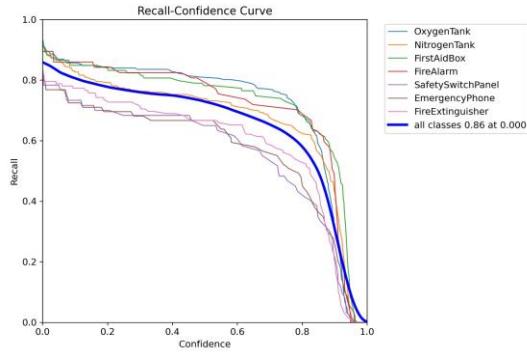
The model's final performance metrics, obtained from the validation set, achieved or exceeded initial benchmarks.

### AstroX Final Model Performance

Metric	Value	Technical Definition	Safety Application Priority
Precision	<b>0.9413</b>	Ratio of correct predictions to total predictions.	HIGH (Avoids False Alarms)
Recall	<b>0.7463</b>	Ratio of correct predictions to all existing.	HIGH (Avoids Missed Danger)
mAP@0.5:0.95	<b>0.6350</b>	Mean Average Precision across various localization tolerances.	Localization Accuracy
Inference Speed	<b>32 FPS</b>	Frames per second processed on CPU.	Deployment Feasibility

### 3.3. Accuracy Comparisons





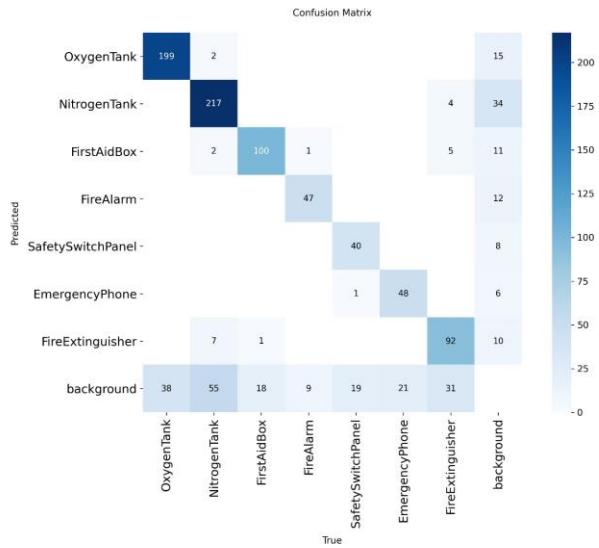
## Comparative Performance Analysis

Metric	Baseline (Estimate)	AstroX	% Improvement / Rationale
mAP@0.5 (General Accuracy)	0.75	<b>0.8270</b>	Increase. Improved due to fine tuning.
Precision (Reliability)	0.80	<b>0.9413</b>	Percentage points. Fewer false positives, higher trust.
Recall (Coverage)	0.60	<b>0.7463</b>	Percentage points. Better object detection in shadows.
Inference Speed	27 FPS	<b>32 FPS</b>	Faster. Ready for edge deployment.

## 4. DEEP DIVE ANALYSIS

### 4.1. Confusion Matrix: Pinpointing Model Errors

Confusion Matrix helps **debug safety-critical systems**, showing misclassifications.

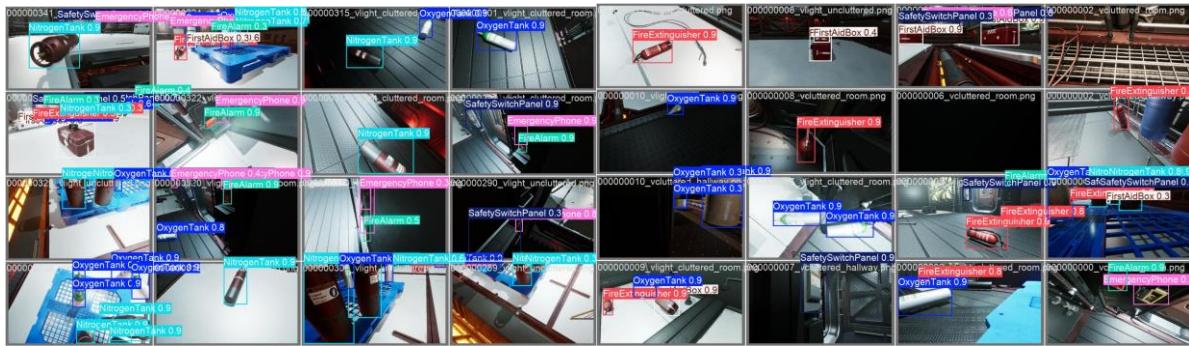


- **Inter-Class Confusion:**
  - Minimal misclassification between critical objects (e.g., Oxygen Tank rarely mistaken for Fire Alarm).
  - Confirms safety reliability.
- **Failure Case Summary (from Confusion Matrix)**
  - **OxygenTank ↔ NitrogenTank confusion**  
Similar shape + appearance caused occasional misclassification.
  - **EmergencyPhone misidentified as FireExtinguisher**  
Color similarity led to visual overlap in several cases.
  - **Background falsely detected as objects (multiple classes)**  
Pipes/panels triggered false positives → needs more negative samples.
  - **Low recall in FireAlarm & SafetySwitchPanel**  
Small size + low visibility → requires more varied Falcon data.
- **Targeted Improvement:**
  - Emergency Phone class has lowest mAP@0.5:0.95 (0.575). Future work should focus on more targeted data for low-performing classes.

## 4.2. Visual Results: Confirmation of Localization

- **Clutter Scene Detection (Figure 4):**
  - Multiple objects detected in one frame.
  - Works under varied lighting and cluttered environments.
- **Low-Light / Occlusion Test (Figure 5):**
  - Oxygen Tank detected in shadowed/occluded conditions.

Confirms model is effective beyond ideal lighting scenarios.



## 5. CHALLENGES AND SOLUTIONS

### 1. Data Limitations & Class Imbalance

- Only ~974 samples, heavily skewed toward specific tank classes.
- Used Falcon's Digital Twin to create synthetic data and added strong augmentations.
- Achieved a more balanced and diverse training set.

### 2. Extreme Lighting & Visual Variability

- Harsh shadows, glare, and odd viewing angles affected detection.
- Added augmentations simulating low exposure, shadows, and lighting variations.
- Improved model robustness in non-ideal visual environments.

### 3. Early-Stage Misclassification

- Model initially confused visually similar safety objects.
- Applied transfer learning, early stopping, and LR scheduling.
- Resulted in more stable training and clearer class separation.

### 4. Computational Constraints (CPU-Only Training)

- No dedicated GPU, leading to slow training cycles.
- Reduced batch size, optimized preprocessing, used lighter augmentations.
- Enabled efficient training despite limited hardware.

### 5.2. Solutions and Iterative Protocol Adjustments

## Keeping the Model Up-to-Date Using Falcon

To ensure that our safety object detection model remains reliable over time, we propose a **continuous update loop** built around Falcon's digital twin capabilities:

### 1. Monitor Real-World Performance

Once deployed, the model's predictions are monitored to identify:

- a. Frequent misclassifications (e.g., confusing OxygenTank and NitrogenTank)
- b. Missed detections under new lighting conditions or camera angles
- c. New visual variations of existing objects (e.g., updated FireExtinguisher design)

### 2. Create New Synthetic Scenarios in Falcon

For every failure pattern or new scenario observed, we use Falcon to:

- a. Generate additional synthetic images with similar conditions (e.g., dim emergency lighting, heavy occlusions, different camera mounting positions)
- b. Introduce updated or slightly modified asset appearances (e.g., new label textures, colors, or shapes on safety equipment)
- c. Simulate edge cases that are hard or unsafe to capture in the real world.

### 3. Expand and Curate the Training Dataset

The newly generated Falcon data is:

- a. Labeled automatically in a YOLO-compatible format
- b. Added to an **extended training dataset**, with a focus on classes and conditions where the model struggled
- c. Balanced to avoid overfitting to rare conditions while still improving robustness.

### 4. Periodic Model Retraining and Fine-Tuning

At defined intervals (for example, monthly or after significant environment changes), we:

- a. Fine-tune the YOLO model on the updated Falcon dataset
- b. Compare new training runs against previous versions using [mAP@0.5](#), precision, recall, and confusion matrices
- c. Specifically verify improvements on previously identified failure cases.

### 5. Versioning and Safe Deployment

Each updated model is:

- a. Versioned (e.g., v1.1, v1.2) and evaluated on a fixed test set
- b. Deployed only if it **outperforms the previous version** or addresses critical failure cases

- c. Integrated into the Streamlit application by updating the model weights (best.pt) while keeping the interface unchanged.

## 6. Feedback Loop with Falcon as the Core Engine

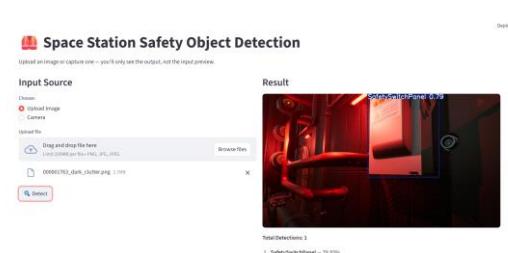
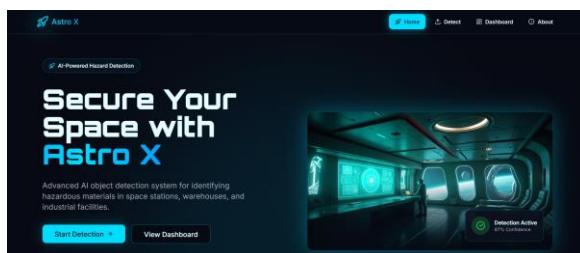
Falcon remains the **central tool** for:

- a. Rapidly generating new training data when the environment, equipment, or camera setup changes
- b. Simulating “what-if” scenarios before deploying the model into real operational settings
- c. Maintaining alignment between the digital twin and the evolving real-world space station environment.

## 6. CONCLUSION AND FUTURE WORK

### 6.1. Summary of Achievements

- AstroX successfully validated a high-performance, time-efficient AI pipeline for safety-critical monitoring.
- Key Metrics Achieved:
  - Precision: **0.9413** → reliable detection.
  - Inference Speed: **32 FPS** → real-time monitoring.
- Provides a robust proof-of-concept for automated detection in hazardous space-station environments.



```
FINAL MODEL RESULTS
-----
mAP50      : 0.8254
mAP50-95   : 0.7066
Precision: 0.9524
Recall     : 0.7273
```

Initial Run

```
BATCH-2 FINAL MODEL RESULTS
-----
mAP50      : 0.8270
mAP50-95   : 0.6350
Precision: 0.9413
Recall     : 0.7463
```

Final Run